

CS305

Computer Architecture

Virtual Memory

Bhaskaran Raman
Room 406, KR Building
Department of CSE, IIT Bombay

<http://www.cse.iitb.ac.in/~br>

Programmer's Burden: Program Size > Main Memory Size



Initially Part1, Part2



Then Part1, Part3



Then Part1, Part4

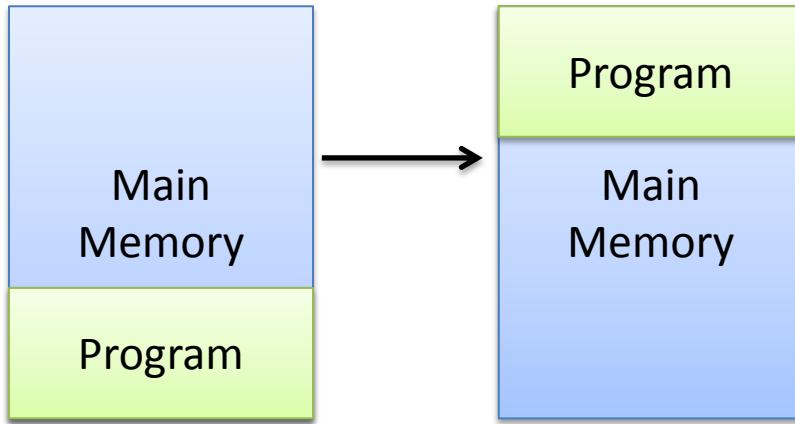


Program had to deal with bringing in relevant parts of the program, into main memory!

Not a problem today: most programs < memory size

Illusion required: memory larger than main memory

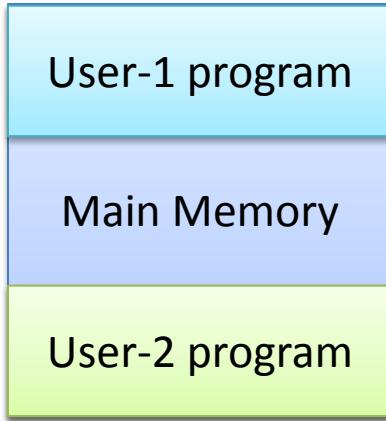
Loader's Burden: Program Relocation



Program, or parts of program need to be constantly rewritten, if loaded to different parts of memory

Illusion required: program at same memory location

OS's Burden: Program Safety



Different user's programs must be protected from one another: malicious or buggy behaviour

Illusion required: each program thinks it has the entire memory to itself, it is the only program

Virtual Memory: Illusions Galore

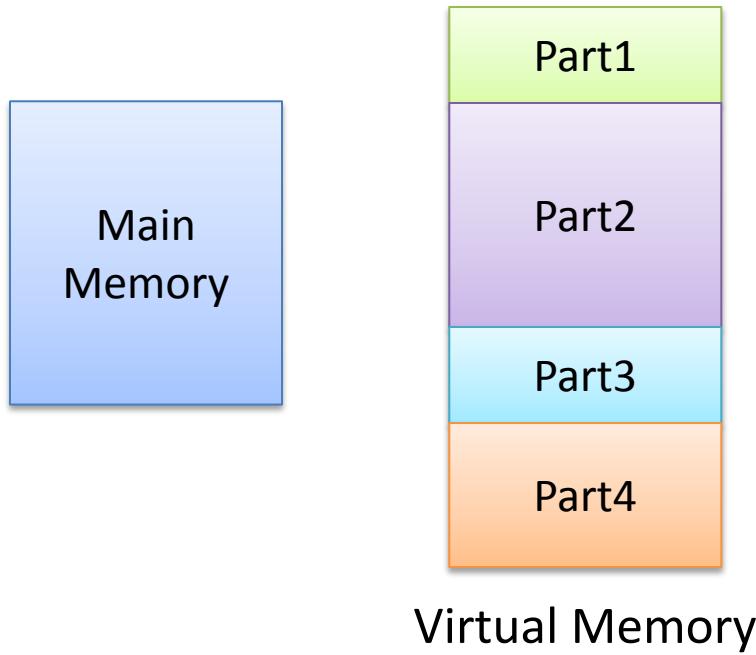
Illusion-1: Larger memory than main memory

Illusion-2: Program at same memory location

Illusion-3: The current program is the only program, with entire memory to itself

Illusion-N... E.g. with shared dynamic libraries, shared memory, etc.

Virtual Memory: The Cache Perspective

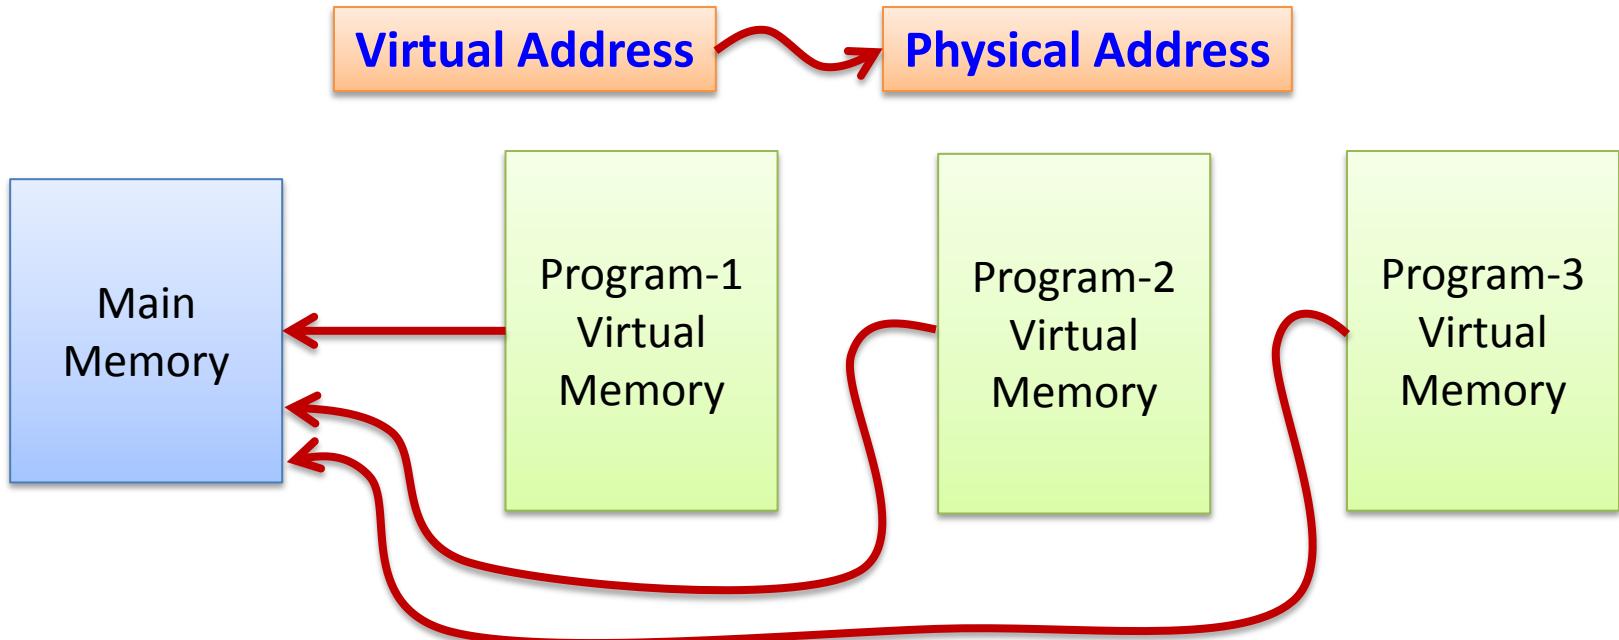


Main memory is a cache for virtual memory

Removes programmer's burden

Incomplete perspective:
e.g. main memory can be
8GB, virtual memory 4GB

Virtual Memory: The Level-of-Indirection Perspective



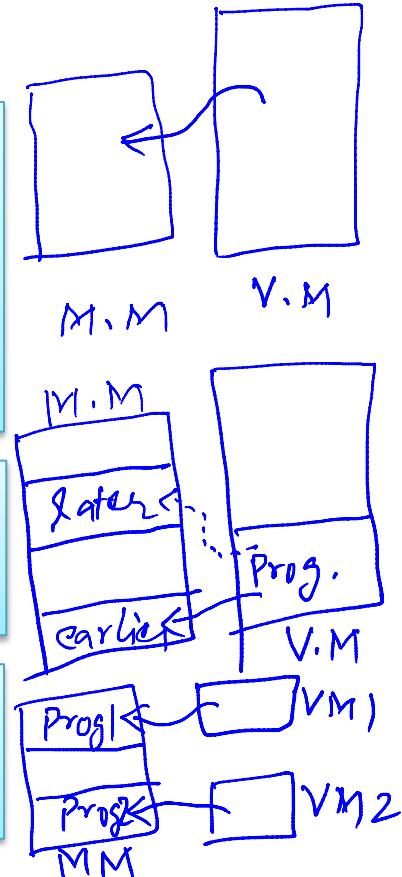
OS, with hardware support, maintains VA → PA mapping

Level-of-Indirection: The Basis for All Illusions

Virtual address (VA) → physical address (PA) mapping:
Like main memory address → cache location mapping
Can support much larger virtual address space than physical (main) memory

Program relocation: change VA → PA mapping!
Program (in terms of VA) remains unchanged

Program safety: separate VA space for each program
No chance of inadvertent access by another program



The Cache Perspective

L1 as Cache for MM

Unit of transfer = block

Not in cache → miss

Mapping: using MM addr, tags

Miss penalty: a few 100 cycles

MM as Cache for VM

Unit of transfer = page

Not in cache → page fault

Mapping: using page table, other schemes

Miss penalty: to disk → a few ms!
Millions or billions of cycles!

VM Design Decisions in the Cache Perspective

Strategy-1: Reduce miss rate

Page size: large (e.g. 32KB)

Associativity: full

Write-back scheme

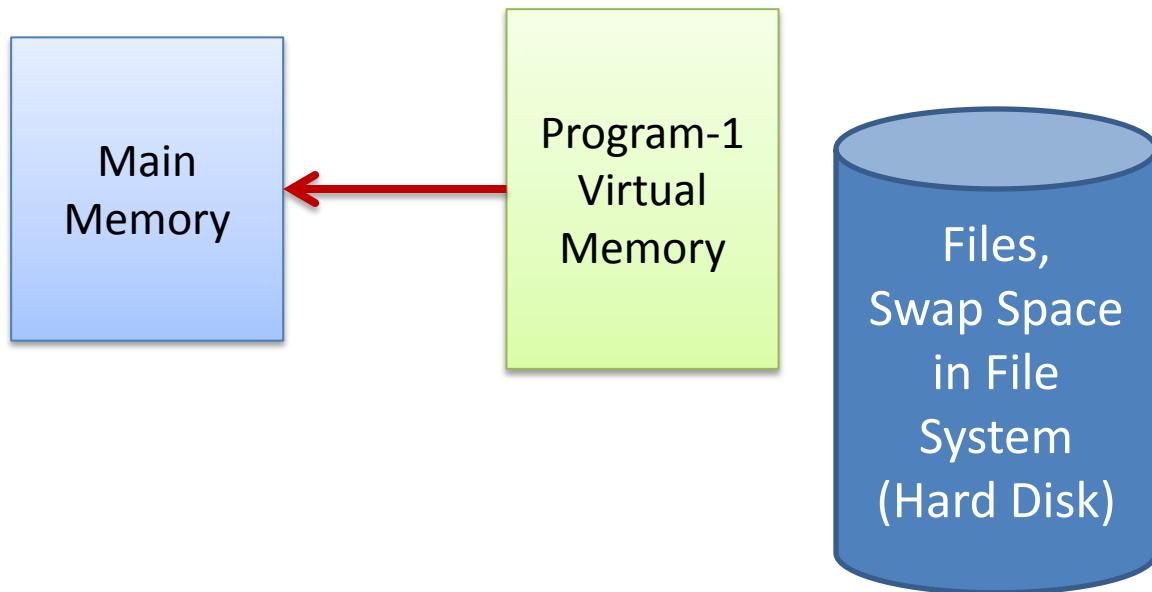
Write-allocate scheme

Page fault: handled in software (OS), sophisticated replacement policy

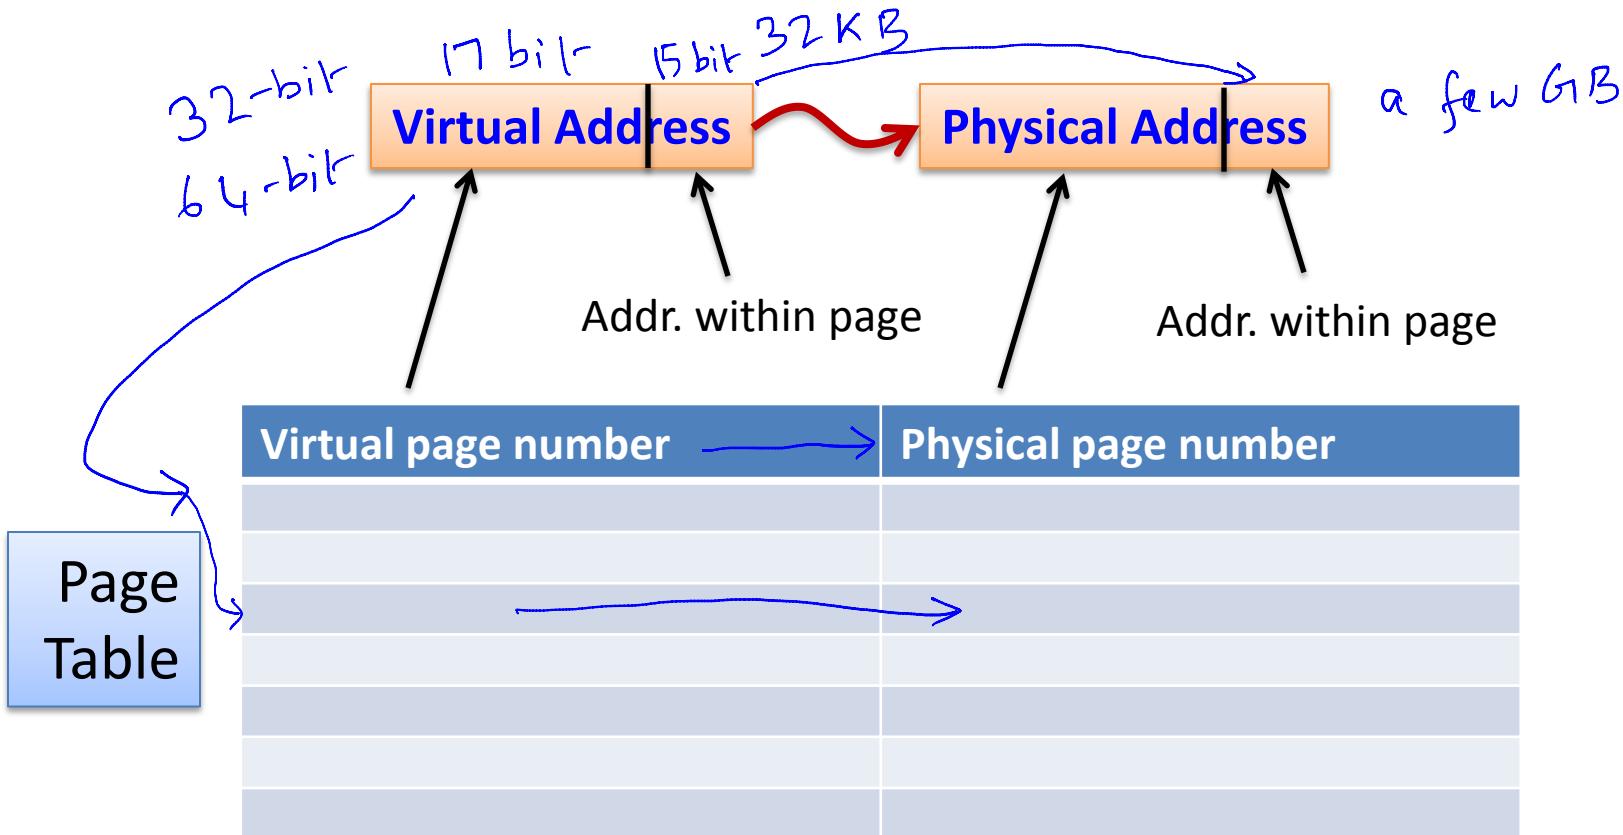
Strategy-2: Multi-tasking

When a program has a miss (page fault), run some other program!

The Swap Space



The VA-to-PA Mapping: Page Table



Page Table: A Numeric Example

32-bit VA space, 8KB pages

Say, 4 bytes per page table entry

Page table size = ?

pages in VM = $2^{32}/2^{13} = 2^{19}$

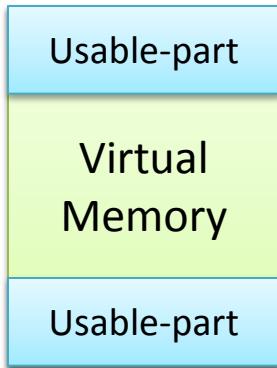
Page table size = 2^{21} bytes = 2MB !

This is per process !

64-bit VA space → page table much larger!

What can we do?

Dealing with Large Mapping Size: Solution-1: Page Table Limits

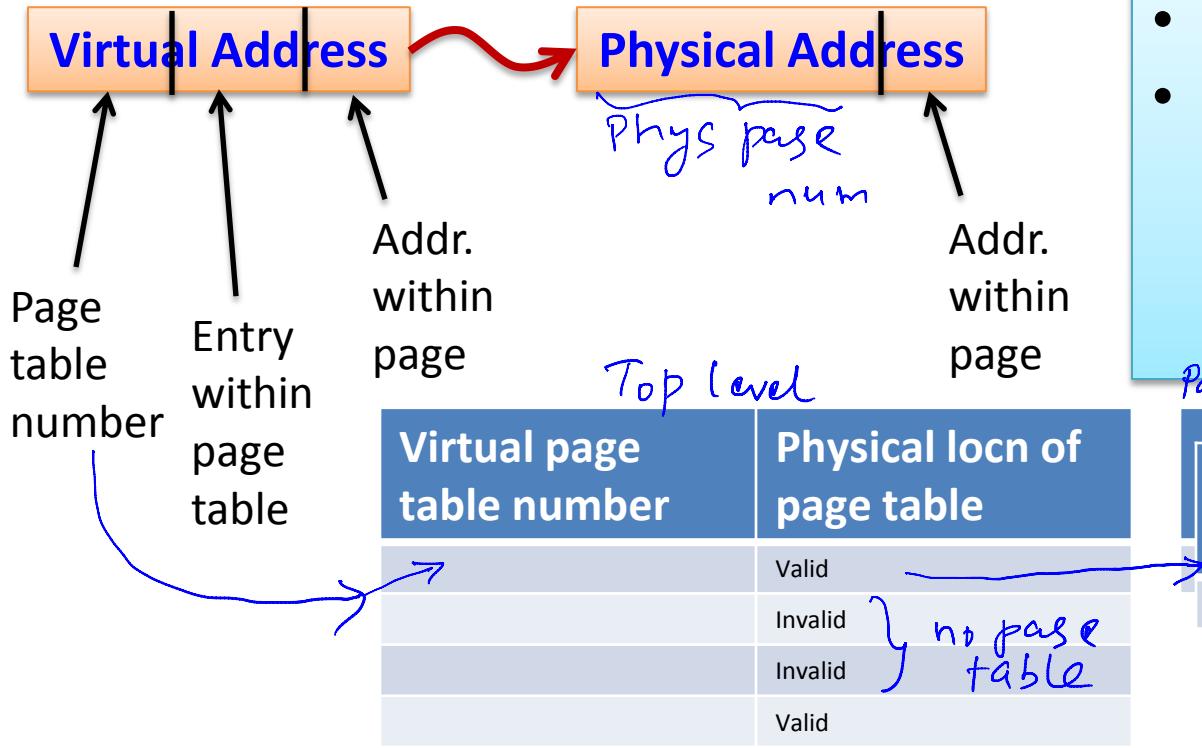


Dealing with Large Mapping Size: Solution-2: Inverted Page Table

Virtual page number	Physical page number

- One entry per physical page
- Entries hashed on virtual page number

Dealing with Large Mapping Size: Solution-3: Multi-level Page Table

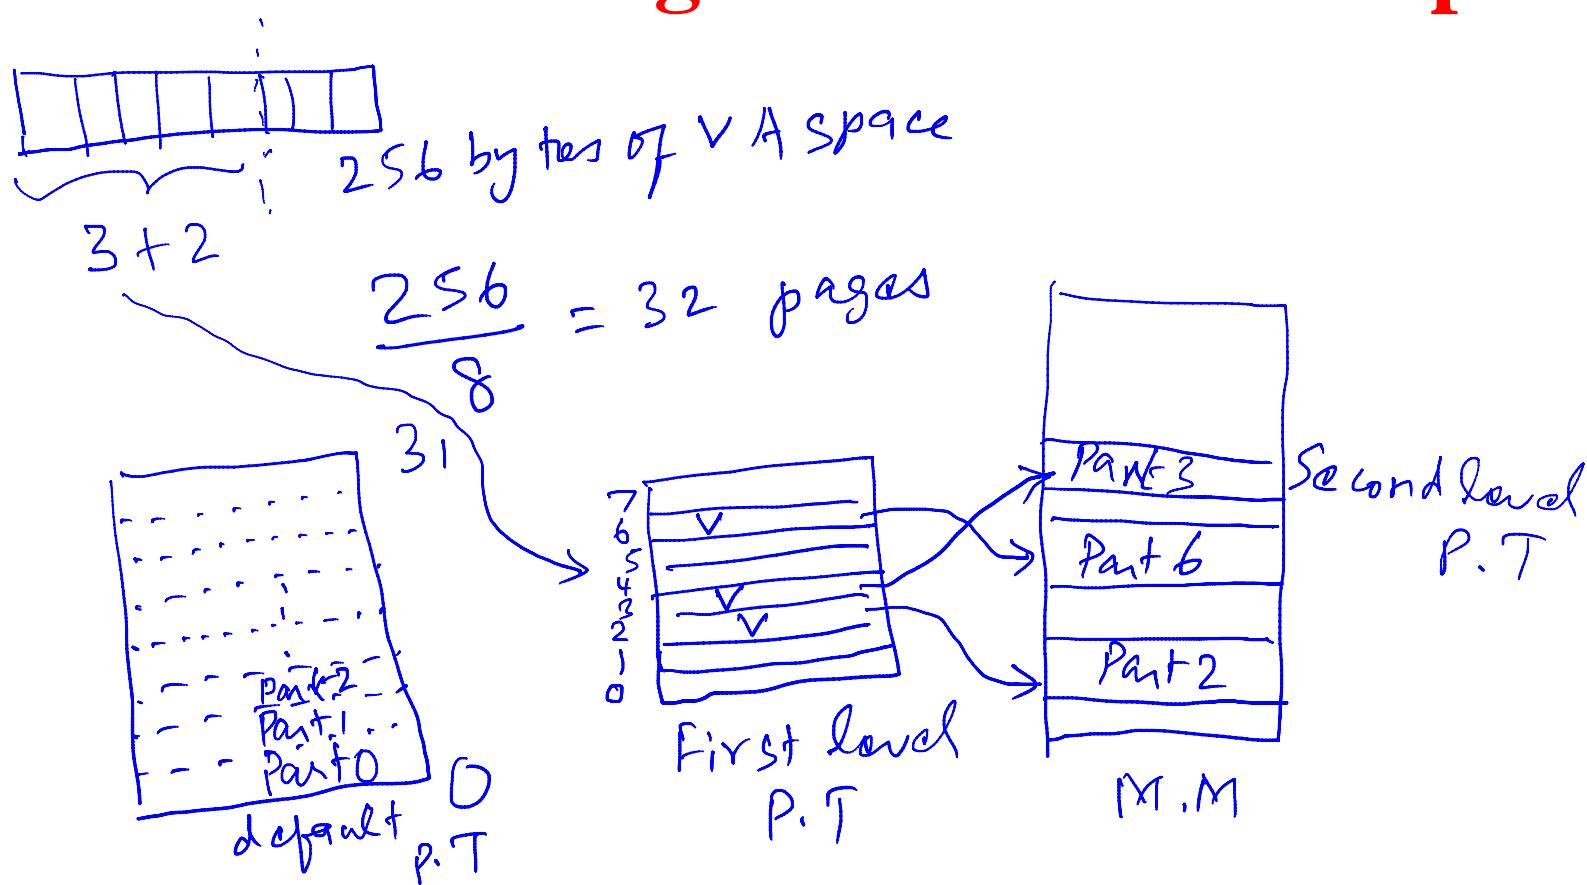


- Main idea: hierarchy
- Idea works since most entries in first level page table will be invalid

Page tables

Page table entry number	Physical page number
Physical address 1	Physical address 1
Physical address 2	Physical address 2
Physical address 3	Physical address 3
Physical address 4	Physical address 4

Multi-Level Page Table: An Example

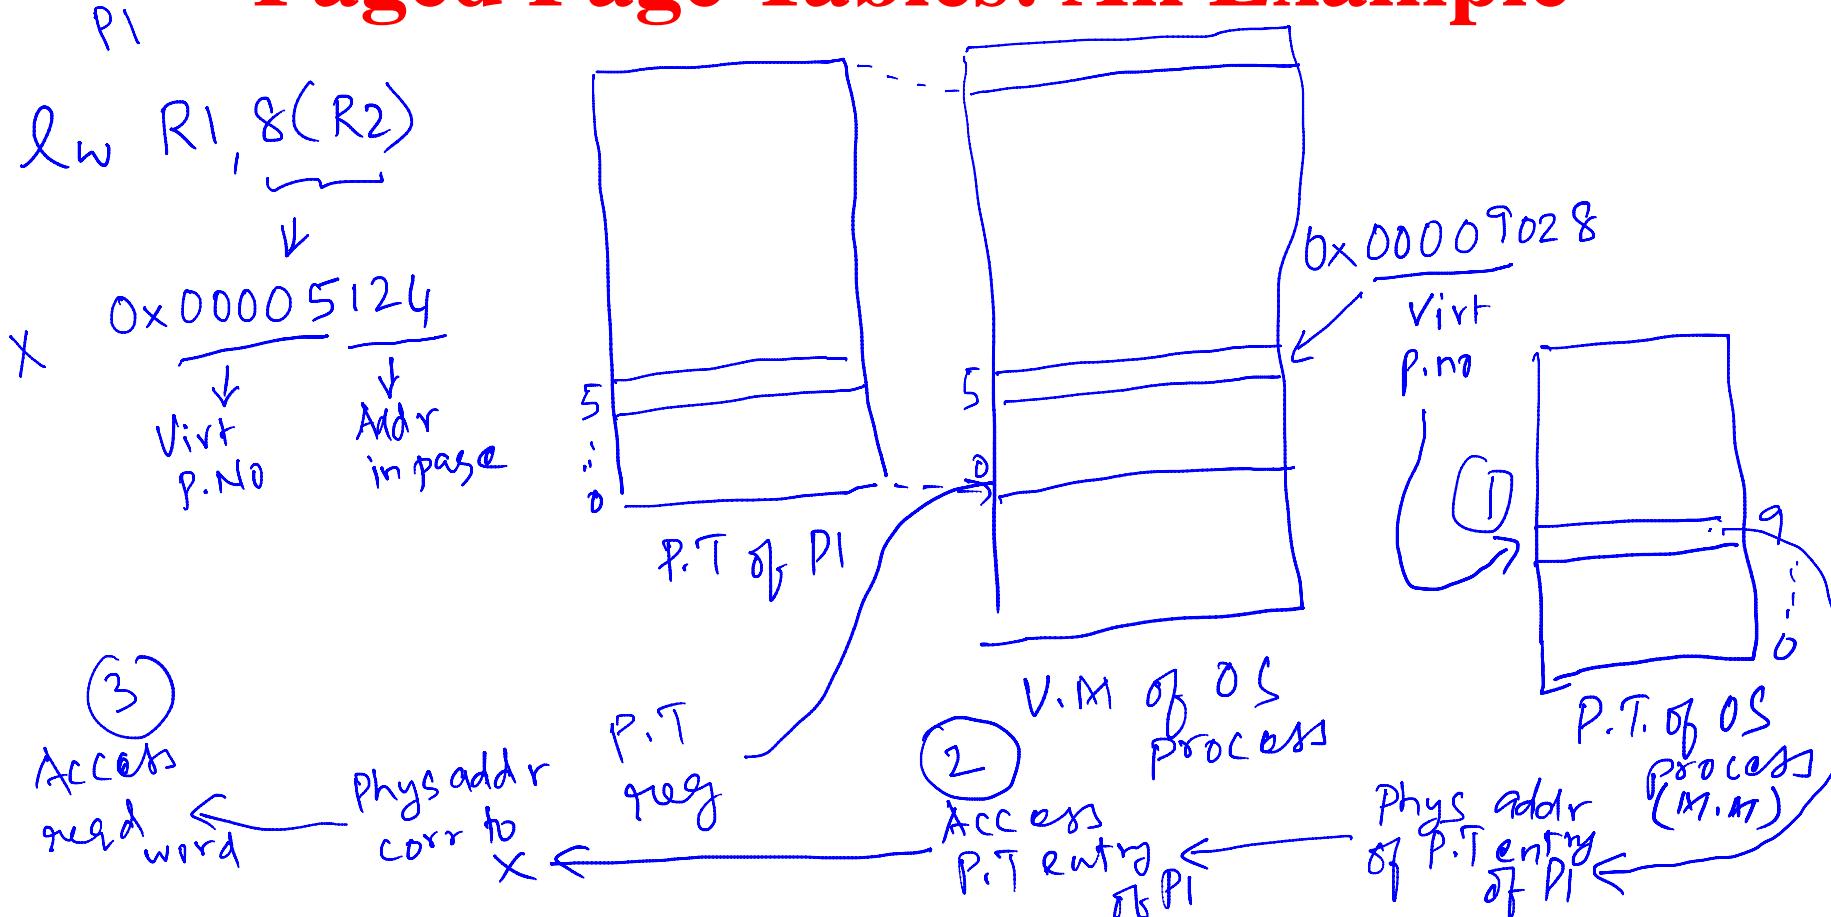


Dealing with Large Mapping Size: Solution-4: Paged Page Tables

Virtual page number	Physical page number

- Main ideas: level of indirection, caching
 - Earlier: page table was in physical memory
 - Now: page table in virtual memory (of a special OS process)
 - Can manage even if page table size is large
 - Page table of OS process pinned to physical memory

Paged Page Tables: An Example



Dealing with Large Mapping Size: Idea Behind the Solutions

- Page table limits: works since most of the page table is empty
- Inverted page table: works since most of the page table is empty
- Multi-level page table: works since most of the page table is empty
- Paged page table: works since there is locality of reference in virtual page references

Summary

- Virtual memory: VA-to-PA mapping is a level of indirection
 - Provides immense flexibility
- Segments: unequal sized blocks
- Managing the mapping: page table
- Ideas to deal with mapping size: page table limits, inverted page table, multi-level page table, paged page table
- Next: putting together VM, caches